

europython
Edinburgh 23-29 July

2018



Reliability of Distributed Systems



- Timeouts
 - Circuit breakers
 - Repeaters
 - Disaster recovery
 - Monitoring
 - Anomaly detection
 - Alerting
 - Error tracking
 - Logging
 - Postmortems
-

Asyncio in production



- Asyncio
 - Async web frameworks
 - Migration to async
 - Async problems
 - Benchmarking
-

Deep Learning with PyTorch for fun and profit



- Style Transfer in Comics
 - External GPUs
 - PyTorch
 - Jupyter
 - Writing Shakespeare
 - Creating audio dramas
-

When to use Machine Learning: Tips, Tricks and Warnings



- Linear regression
 - Battery usage
 - Whereami
 - Car damage recognition
 - Transfer learning
 - Code autocompletion
 - Cryptocurrency trading
-

Industrial Machine Learning Pipelines with Python and Airflow



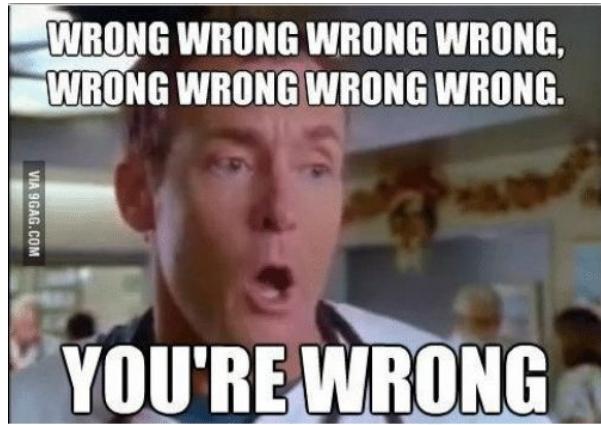
- Airflow
 - Model training
 - Scaling
 - Data transformation
 - DAG
 - Deep RNNs
-

The Web is terrifying! Using PyData stack to spy on spies



- Jupyter
 - Pandas
 - Dask
 - Bokeh
 - Scikit-Learn
 - NetworkX
 - Mozilla .profile
 - Cookie syncing
 - Zombie cookies
 - Identity resolution industry
-

Asyncio in practice: We did it wrong!



- Asyncio
- Concurrency
- PUB/SUB
- Best Practices

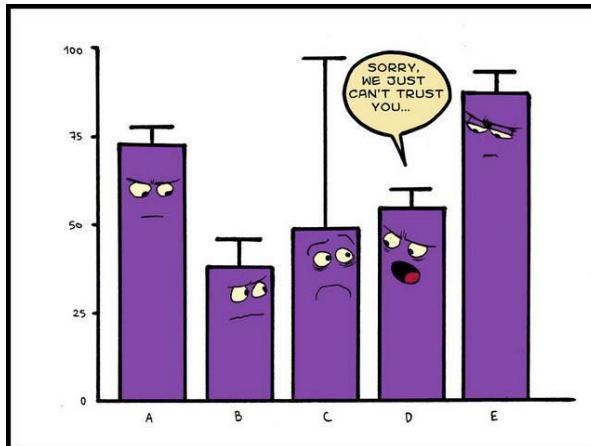
rogue.ly/aio

Productionizing your ML code seamlessly



- by Yelp guy
 - ML system components
 - Success measurement
 - General advice
 - Logging
 - Pipeline monitoring
-

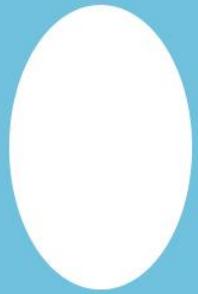
Lies, damned lies, and statistics



- Correlation
- Lurking variables
- Slicing and dicing data
- Bias
- Data visualization
- Statistical significance

MISLEADING STATISTICS

THE AVERAGE ADULT HAS ONE TESTICLE



Meaningful > Just True

CONTENT SHOULD BE USEFUL, NOT JUST PRETTY

vert.ms/Baddata




© marketoonist.com

Having fun with music and Keras



- Building a synthesizer using numpy
 - Generating music
 - Auto Encoders
 - Keras
-







Conclusion

